



## Spark Release 3.3.4

Spark 3.3.4 is the last maintenance release containing security and correctness fixes. This release is based on the branch-3.3 maintenance branch of Spark. We strongly recommend all 3.3 users to upgrade to this stable release.

### Notable changes

- [SPARK-43327] (<https://issues.apache.org/jira/browse/SPARK-43327>): Trigger `committer.setupJob` before plan execute in `FileFormatWriter#write`
- [SPARK-43393] (<https://issues.apache.org/jira/browse/SPARK-43393>): Address sequence expression overflow bug
- [SPARK-44547] (<https://issues.apache.org/jira/browse/SPARK-44547>): Ignore fallback storage for cached RDD migration
- [SPARK-44581] (<https://issues.apache.org/jira/browse/SPARK-44581>): Fix the bug that `ShutdownHookManager` gets wrong UGI from `SecurityManager` of `ApplicationMaster`
- [SPARK-44725] (<https://issues.apache.org/jira/browse/SPARK-44725>): Document `spark.network.timeoutInterval`
- [SPARK-44805] (<https://issues.apache.org/jira/browse/SPARK-44805>): `getBytes/getShorts/getInts/etc.` should work in a column vector that has a dictionary
- [SPARK-44857] (<https://issues.apache.org/jira/browse/SPARK-44857>): Fix `getBaseURI` error in Spark Worker LogPage UI buttons
- [SPARK-44871] (<https://issues.apache.org/jira/browse/SPARK-44871>): Fix `percentile_disc` behaviour
- [SPARK-44920] (<https://issues.apache.org/jira/browse/SPARK-44920>): Use `await()` instead of `awaitUninterruptibly()` in `TransportClientFactory.createClient()`
- [SPARK-44925] (<https://issues.apache.org/jira/browse/SPARK-44925>): K8s default service token file should not be materialized into token
- [SPARK-44935] (<https://issues.apache.org/jira/browse/SPARK-44935>): Fix `RELEASE` file to have the correct information in Docker images if exists
- [SPARK-44937] (<https://issues.apache.org/jira/browse/SPARK-44937>): Mark connection as `timedOut` in `TransportClient.close`

- [SPARK-44973] (<https://issues.apache.org/jira/browse/SPARK-44973>): Fix `ArrayIndexOutOfBoundsException` in `conv()`
- [SPARK-44990] (<https://issues.apache.org/jira/browse/SPARK-44990>): Reduce the frequency of `get spark.sql.legacy.nullValueWrittenAsQuotedEmptyStringCsv`
- [SPARK-45057] (<https://issues.apache.org/jira/browse/SPARK-45057>): Avoid acquire read lock when `keepReadLock` is false
- [SPARK-45079] (<https://issues.apache.org/jira/browse/SPARK-45079>): Fix an internal error from `percentile_approx()` on NULL accuracy
- [SPARK-45100] (<https://issues.apache.org/jira/browse/SPARK-45100>): Fix an internal error from `reflect()` on NULL class and method
- [SPARK-45187] (<https://issues.apache.org/jira/browse/SPARK-45187>): Fix `WorkerPage` to use the same pattern for `logPage` urls
- [SPARK-45227] (<https://issues.apache.org/jira/browse/SPARK-45227>): Fix a subtle thread-safety issue with `CoarseGrainedExecutorBackend`
- [SPARK-45389] (<https://issues.apache.org/jira/browse/SPARK-45389>): Correct `MetaException` matching rule on getting partition metadata
- [SPARK-45430] (<https://issues.apache.org/jira/browse/SPARK-45430>): Fix for `FramelessOffsetWindowFunction` when IGNORE NULLS and `offset > rowCount`
- [SPARK-45508] (<https://issues.apache.org/jira/browse/SPARK-45508>): Add “`--add-opens=java.base/jdk.internal.ref=ALL-UNNAMED`” so Platform can access Cleaner on Java 9+
- [SPARK-45580] (<https://issues.apache.org/jira/browse/SPARK-45580>): Handle case where a nested subquery becomes an existence join
- [SPARK-45670] (<https://issues.apache.org/jira/browse/SPARK-45670>): `SparkSubmit` does not support `--total-executor-cores` when deploying on K8s
- [SPARK-45749] (<https://issues.apache.org/jira/browse/SPARK-45749>): Fix Spark History Server to sort `Duration` column properly
- [SPARK-45920] (<https://issues.apache.org/jira/browse/SPARK-45920>): group by ordinal should be idempotent
- [SPARK-46006] (<https://issues.apache.org/jira/browse/SPARK-46006>): `YarnAllocator` miss clean `targetNumExecutorsPerResourceProfileId` after `YarnSchedulerBackend` call stop

- [SPARK-46012] (<https://issues.apache.org/jira/browse/SPARK-46012>): EventLogFileReader should not read rolling logs if app status file is missing
- [SPARK-46029] (<https://issues.apache.org/jira/browse/SPARK-46029>): Escape the single quote, \_ and % for DS V2 pushdown
- [SPARK-46092] (<https://issues.apache.org/jira/browse/SPARK-46092>): Don't push down Parquet row group filters that overflow
- [SPARK-46095] (<https://issues.apache.org/jira/browse/SPARK-46095>): Document REST API for Spark Standalone Cluster
- [SPARK-46239] (<https://issues.apache.org/jira/browse/SPARK-46239>): Hide Jetty info
- [SPARK-46286] (<https://issues.apache.org/jira/browse/SPARK-46286>): Document `spark.io.compression.zstd.bufferPool.enabled`

## Dependency Changes

While being a maintenance release we did still upgrade some dependencies in this release they are:

- [SPARK-45885] (<https://issues.apache.org/jira/browse/SPARK-45885>): Upgrade ORC to 1.7.10

You can consult JIRA for the detailed changes (<https://s.apache.org/spark-3.3.4>).

We would like to acknowledge all community members for contributing patches to this release.

Spark News Archive (</news/>)

### Latest News

Spark 3.5.5 released (</news/spark-3-5-5-released.html>) (Feb 27, 2025)

Spark 3.5.4 released (</news/spark-3-5-4-released.html>) (Dec 20, 2024)

Spark 3.4.4 released (</news/spark-3-4-4-released.html>) (Oct 27, 2024)

Preview release of Spark 4.0 (</news/spark-4.0.0-preview2.html>) (Sep 26, 2024)

[Archive \(/news/index.html\)](/news/index.html)



(<https://www.apache.org/events/current-event.html>)

## **DOWNLOAD SPARK (/DOWNLOADS.HTML)**

Built-in Libraries:

SQL and DataFrames ([/sql/](#))

Spark Streaming ([/streaming/](#))

MLlib (machine learning) ([/mllib/](#))

GraphX (graph) ([/graphx/](#))

Third-Party Projects ([/third-party-projects.html](#))

Apache Spark, Spark, Apache, the Apache feather logo, and the Apache Spark project logo are either registered trademarks or trademarks of The Apache Software Foundation in the United States and other countries. See guidance on use of Apache Spark trademarks ([/trademarks.html](#)). All other marks mentioned may be trademarks or registered trademarks of their respective owners. Copyright © 2018 The Apache Software Foundation, Licensed under the Apache License, Version 2.0 (<https://www.apache.org/licenses/>).