



Spark Release 3.4.4

Spark 3.4.4 is the last maintenance release containing security and correctness fixes. This release is based on the branch-3.4 maintenance branch of Spark. We strongly recommend all 3.4 users to upgrade to this stable release.

Notable changes

- [SPARK-43242] (<https://issues.apache.org/jira/browse/SPARK-43242>): Fix throw 'Unexpected type of BlockId' in shuffle corruption diagnose
- [SPARK-45988] (<https://issues.apache.org/jira/browse/SPARK-45988>): Fix typehints to handle List GenericAlias in Python 3.11+
- [SPARK-46535] (<https://issues.apache.org/jira/browse/SPARK-46535>): Fix NPE when describe extended a column without col stats
- [SPARK-46957] (<https://issues.apache.org/jira/browse/SPARK-46957>): Decommission migrated shuffle files should be able to cleanup from executor
- [SPARK-47129] (<https://issues.apache.org/jira/browse/SPARK-47129>): Make ResolveRelations cache connect plan properly
- [SPARK-47172] (<https://issues.apache.org/jira/browse/SPARK-47172>): Add support for AES-GCM for RPC encryption
- [SPARK-47828] (<https://issues.apache.org/jira/browse/SPARK-47828>): DataFrameWriterV2.overwrite fails with invalid plan
- [SPARK-47895] (<https://issues.apache.org/jira/browse/SPARK-47895>): group by all should be idempotent
- [SPARK-47897] (<https://issues.apache.org/jira/browse/SPARK-47897>): Fix ExpressionSet performance regression in scala 2.12
- [SPARK-47927] (<https://issues.apache.org/jira/browse/SPARK-47927>): Fix nullability attribute in UDF decoder
- [SPARK-48016] (<https://issues.apache.org/jira/browse/SPARK-48016>): Fix a bug in try_divide function when with decimals
- [SPARK-48019] (<https://issues.apache.org/jira/browse/SPARK-48019>): Fix incorrect behavior in ColumnVector/ColumnarArray with dictionary and nulls
- [SPARK-48037] (<https://issues.apache.org/jira/browse/SPARK-48037>): Fix SortShuffleWriter lacks shuffle write related metrics resulting in potentially inaccurate data

- [SPARK-48081] (<https://issues.apache.org/jira/browse/SPARK-48081>): Fix ClassCastException in NTile.checkInputDataTypes() when argument is non-foldable or of wrong type
- [SPARK-48105] (<https://issues.apache.org/jira/browse/SPARK-48105>): Fix the race condition between state store unloading and snapshotting
- [SPARK-48128] (<https://issues.apache.org/jira/browse/SPARK-48128>): For BitwiseCount / bit_count expression, fix codegen syntax error for boolean type inputs
- [SPARK-48155] (<https://issues.apache.org/jira/browse/SPARK-48155>): AQEPropagateEmptyRelation for join should check if remain child is just BroadcastQueryStageExec
- [SPARK-48172] (<https://issues.apache.org/jira/browse/SPARK-48172>): Fix escaping issues in JDBC Dialects
- [SPARK-48248] (<https://issues.apache.org/jira/browse/SPARK-48248>): Fix nested array to respect legacy conf of inferArrayTypeFromFirstElement
- [SPARK-48292] (<https://issues.apache.org/jira/browse/SPARK-48292>): Revert SPARK-39195 Spark OutputCommitCoordinator should abort stage when committed file not consistent with task status
- [SPARK-48484] (<https://issues.apache.org/jira/browse/SPARK-48484>): Fix: V2Write use the same TaskAttemptId for different task attempts
- [SPARK-48642] (<https://issues.apache.org/jira/browse/SPARK-48642>): False SparkOutOfMemoryError caused by killing task on spilling
- [SPARK-48710] (<https://issues.apache.org/jira/browse/SPARK-48710>): Limit NumPy version to supported range ($\geq 1.15, < 2$)
- [SPARK-48759] (<https://issues.apache.org/jira/browse/SPARK-48759>): Add migration doc for CREATE TABLE AS SELECT behavior change since Spark 3.4
- [SPARK-48791] (<https://issues.apache.org/jira/browse/SPARK-48791>): Fix perf regression caused by the accumulators registration overhead using CopyOnWriteArrayList
- [SPARK-48930] (<https://issues.apache.org/jira/browse/SPARK-48930>): Redact awsAccessKeyId by including accesskey pattern
- [SPARK-48934] (<https://issues.apache.org/jira/browse/SPARK-48934>): Python datetime types converted incorrectly for setting timeout in applyInPandasWithState
- [SPARK-48965] (<https://issues.apache.org/jira/browse/SPARK-48965>): Use the correct schema in Dataset#toJSON

- [SPARK-48991] (<https://issues.apache.org/jira/browse/SPARK-48991>): Move path initialization into try-catch block in FileStreamSink.hasMetadata
- [SPARK-49000] (<https://issues.apache.org/jira/browse/SPARK-49000>): Fix “select count(distinct 1) from t” where t is empty table by expanding RewriteDistinctAggregates
- [SPARK-49005] (<https://issues.apache.org/jira/browse/SPARK-49005>): Use 17-j ammy tag instead of 17-j re to prevent Python 3.12
- [SPARK-49038] (<https://issues.apache.org/jira/browse/SPARK-49038>): SQLMetric should report the raw value in the accumulator update event
- [SPARK-49039] (<https://issues.apache.org/jira/browse/SPARK-49039>): Reset checkbox when executor metrics are loaded in the Stages tab
- [SPARK-49094] (<https://issues.apache.org/jira/browse/SPARK-49094>): Fix ignoreCorruptFiles non-functioning for hive orc impl with mergeSchema off
- [SPARK-49176] (<https://issues.apache.org/jira/browse/SPARK-49176>): Fix spark.ui.custom.executor.log.url docs by adding K8s
- [SPARK-49179] (<https://issues.apache.org/jira/browse/SPARK-49179>): Fix v2 multi bucketed inner joins throw AssertionError
- [SPARK-49182] (<https://issues.apache.org/jira/browse/SPARK-49182>): Stop publish site/docs/{version}/api/python/_sources dir
- [SPARK-49193] (<https://issues.apache.org/jira/browse/SPARK-49193>): Improve the performance of RowSetUtils.toColumnBasedSet
- [SPARK-49197] (<https://issues.apache.org/jira/browse/SPARK-49197>): Redact Spark Command output in launcher module
- [SPARK-49261] (<https://issues.apache.org/jira/browse/SPARK-49261>): Don't replace literals in aggregate expressions with group-by expressions
- [SPARK-49352] (<https://issues.apache.org/jira/browse/SPARK-49352>): Avoid redundant array transform for identical expression
- [SPARK-49385] (<https://issues.apache.org/jira/browse/SPARK-49385>): Fix getReusablePVCs to use podCreationTimeout instead of podAllocationDelay
- [SPARK-49408] (<https://issues.apache.org/jira/browse/SPARK-49408>): Use IndexedSeq in ProjectingInternalRow
- [SPARK-49595] (<https://issues.apache.org/jira/browse/SPARK-49595>): Fix DataFrame.unpivot/melt in Spark Connect Scala Client

- [SPARK-49628] (<https://issues.apache.org/jira/browse/SPARK-49628>): ConstantFolding should copy stateful expression before evaluating
- [SPARK-49750] (<https://issues.apache.org/jira/browse/SPARK-49750>): Mention delegation token support in K8s mode
- [SPARK-49760] (<https://issues.apache.org/jira/browse/SPARK-49760>): Correct handling of SPARK_USER env variable override in app master
- [SPARK-49804] (<https://issues.apache.org/jira/browse/SPARK-49804>): Fix to use the exit code of executor container always
- [SPARK-49836] (<https://issues.apache.org/jira/browse/SPARK-49836>): Fix possibly broken query when window is provided to window/session_window fn
- [SPARK-49843] (<https://issues.apache.org/jira/browse/SPARK-49843>): Fix change comment on char/varchar columns
- [SPARK-49959] (<https://issues.apache.org/jira/browse/SPARK-49959>): Fix ColumnarArray.copy() to read nulls from the correct offset
- [SPARK-50021] (<https://issues.apache.org/jira/browse/SPARK-50021>): Fix ApplicationPage to hide App UI links when UI is disabled
- [SPARK-50022] (<https://issues.apache.org/jira/browse/SPARK-50022>): Fix MasterPage to hide App UI links when UI is disabled

Dependency Changes

While being a maintenance release we did still upgrade some dependencies in this release they are:

- [SPARK-43394] (<https://issues.apache.org/jira/browse/SPARK-43394>): Upgrade maven to 3.8.8
- [SPARK-45590] (<https://issues.apache.org/jira/browse/SPARK-45590>): Upgrade okio to 1.17.6 from 1.15.0

You can consult JIRA for the detailed changes (<https://s.apache.org/spark-3.4.4>).

We would like to acknowledge all community members for contributing patches to this release.

Spark News Archive (</news/>)

Latest News

Spark 3.5.5 released (</news/spark-3-5-5-released.html>) (Feb 27, 2025)

Spark 3.5.4 released (</news/spark-3-5-4-released.html>) (Dec 20, 2024)

[Spark 3.4.4 released \(/news/spark-3-4-4-released.html\)](/news/spark-3-4-4-released.html) (Oct 27, 2024)

[Preview release of Spark 4.0 \(/news/spark-4.0.0-preview2.html\)](/news/spark-4.0.0-preview2.html) (Sep 26, 2024)

[Archive \(/news/index.html\)](/news/index.html)



[\(/https://www.apache.org/events/current-event.html\)](https://www.apache.org/events/current-event.html)

DOWNLOAD SPARK (/DOWNLOADS.HTML)

Built-in Libraries:

[SQL and DataFrames \(/sql/\)](/sql/)

[Spark Streaming \(/streaming/\)](/streaming/)

[MLlib \(machine learning\) \(/mllib/\)](/mllib/)

[GraphX \(graph\) \(/graphx/\)](/graphx/)

[Third-Party Projects \(/third-party-projects.html\)](/third-party-projects.html)

Apache Spark, Spark, Apache, the Apache feather logo, and the Apache Spark project logo are either registered trademarks or trademarks of The Apache Software Foundation in the United States and other countries. See guidance on use of Apache Spark trademarks (</trademarks.html>). All other marks mentioned may be trademarks or registered trademarks of their respective owners. Copyright © 2018 The Apache Software Foundation, Licensed under the Apache License, Version 2.0 (<https://www.apache.org/licenses/>).