



## Spark Release 3.5.3

Spark 3.5.3 is the third maintenance release containing security and correctness fixes. This release is based on the branch-3.5 maintenance branch of Spark. We strongly recommend all 3.5 users to upgrade to this stable release.

### Notable changes

- [SPARK-49176] (<https://issues.apache.org/jira/browse/SPARK-49176>) - Fix `spark.ui.custom.executor.log.url` docs by adding K8s
- [SPARK-49182] (<https://issues.apache.org/jira/browse/SPARK-49182>) - Stop publish site/docs/{version}/api/python/\_sources
- [SPARK-49197] (<https://issues.apache.org/jira/browse/SPARK-49197>) - Redact Spark Command output in launcher module
- [SPARK-43242] (<https://issues.apache.org/jira/browse/SPARK-43242>) - `diagnoseCorruption` should not throw `Unexpected type of BlockId` for `ShuffleBlockBatchId`
- [SPARK-46037] (<https://issues.apache.org/jira/browse/SPARK-46037>) - When Left Join build Left, `ShuffledHashJoinExec` may result in incorrect results
- [SPARK-46632] (<https://issues.apache.org/jira/browse/SPARK-46632>) - `EquivalentExpressions` throw `IllegalStateException`
- [SPARK-48965] (<https://issues.apache.org/jira/browse/SPARK-48965>) - `toJSON` produces wrong values if `DecimalType` information is lost in `as[Product]`
- [SPARK-49038] (<https://issues.apache.org/jira/browse/SPARK-49038>) - Spark UI is showing wrong timing metrics for min and median time.
- [SPARK-49039] (<https://issues.apache.org/jira/browse/SPARK-49039>) - Reset checkbox when executor metrics are loaded in the Stages tab
- [SPARK-49179] (<https://issues.apache.org/jira/browse/SPARK-49179>) - Fix v2 multi bucketed inner joins throw `AssertionError`
- [SPARK-49183] (<https://issues.apache.org/jira/browse/SPARK-49183>) - `V2SessionCatalog.createTable` should respect `PROP_IS_MANAGED_LOCATION`

- [SPARK-49193] (<https://issues.apache.org/jira/browse/SPARK-49193>) - Improve the performance of `RowSetUtils.toColumnBasedSet`
- [SPARK-49200] (<https://issues.apache.org/jira/browse/SPARK-49200>) - Fix null type non-codegen ordering exception
- [SPARK-49300] (<https://issues.apache.org/jira/browse/SPARK-49300>) - Fix Hadoop delegation token leak when `tokenRenewalInterval` is not set.
- [SPARK-49333] (<https://issues.apache.org/jira/browse/SPARK-49333>) - Shutdown timeout thread while cleaning up `SparkExecuteStatementOperation`
- [SPARK-49352] (<https://issues.apache.org/jira/browse/SPARK-49352>) - Avoid redundant array transform for identical expression
- [SPARK-49359] (<https://issues.apache.org/jira/browse/SPARK-49359>) - allow `StagedTableCatalog` implementations to fall back to non-atomic write
- [SPARK-49381] (<https://issues.apache.org/jira/browse/SPARK-49381>) - Fix SPARK-41388 test case to assert the result
- [SPARK-49385] (<https://issues.apache.org/jira/browse/SPARK-49385>) - Fix `getReusablePVCs` to use `podCreationTimeout` instead of `podAllocationDelay`
- [SPARK-49402] (<https://issues.apache.org/jira/browse/SPARK-49402>) - Fix Binder integration in PySpark documentation
- [SPARK-49476] (<https://issues.apache.org/jira/browse/SPARK-49476>) - Base64 nullability change
- [SPARK-49480] (<https://issues.apache.org/jira/browse/SPARK-49480>) - `NullPointerException` from `SparkThrowableHelper.isInternalError` method
- [SPARK-49509] (<https://issues.apache.org/jira/browse/SPARK-49509>) - Use `Platform.allocateDirectBuffer` instead of `ByteBuffer.allocateDirect`
- [SPARK-49152] (<https://issues.apache.org/jira/browse/SPARK-49152>) - `V2SessionCatalog` should use `V2Command`
- [SPARK-49211] (<https://issues.apache.org/jira/browse/SPARK-49211>) - V2 Catalog can also support built-in data sources
- [SPARK-49246] (<https://issues.apache.org/jira/browse/SPARK-49246>) - `TableCatalog#loadTable` should indicate if it's for writing
- [SPARK-49275] (<https://issues.apache.org/jira/browse/SPARK-49275>) - Fix return type nullness of the xpath expression
- [SPARK-49408] (<https://issues.apache.org/jira/browse/SPARK-49408>) - Poor performance in `ProjectingInternalRow`

# Dependency changes

- [SPARK-45590] (<https://issues.apache.org/jira/browse/SPARK-45590>) - okio-1.15.0 CVE-2023-3635

You can consult JIRA for the detailed changes (<https://s.apache.org/spark-3.5.3>).

We would like to acknowledge all community members for contributing patches to this release.

Spark News Archive (</news/>)

## Latest News

Spark 3.5.5 released (</news/spark-3-5-5-released.html>) (Feb 27, 2025)

Spark 3.5.4 released (</news/spark-3-5-4-released.html>) (Dec 20, 2024)

Spark 3.4.4 released (</news/spark-3-4-4-released.html>) (Oct 27, 2024)

Preview release of Spark 4.0 (</news/spark-4.0.0-preview2.html>) (Sep 26, 2024)

Archive (</news/index.html>)



(<https://www.apache.org/events/current-event.html>)

## DOWNLOAD SPARK ([/DOWNLOADS.HTML](/downloads.html))

Built-in Libraries:

SQL and DataFrames (</sql/>)

Spark Streaming (</streaming/>)

MLlib (machine learning) (</mllib/>)

GraphX (graph) (</graphx/>)

Third-Party Projects (</third-party-projects.html>)

Apache Spark, Spark, Apache, the Apache feather logo, and the Apache Spark project logo are either registered trademarks or trademarks of The Apache Software Foundation in the United States and other countries. See guidance on use of Apache Spark trademarks (</trademarks.html>). All other marks mentioned may be trademarks or registered trademarks of their respective owners. Copyright © 2018 The Apache

Software Foundation, Licensed under the Apache License, Version 2.0  
(<https://www.apache.org/licenses/>).