



Spark Release 3.5.4

Spark 3.5.4 is the third maintenance release containing security and correctness fixes. This release is based on the branch-3.5 maintenance branch of Spark. We strongly recommend all 3.5 users to upgrade to this stable release.

Notable changes

- [SPARK-47702] (<https://issues.apache.org/jira/browse/SPARK-47702>): Remove Shuffle service endpoint from the locations list when RDD block is removed from a node
- [SPARK-48155] (<https://issues.apache.org/jira/browse/SPARK-48155>): AQEPropagateEmptyRelation for join should check if remain child is just BroadcastQueryStageExec
- [SPARK-49261] (<https://issues.apache.org/jira/browse/SPARK-49261>): Don't replace literals in aggregate expressions with group-by expressions
- [SPARK-49294] (<https://issues.apache.org/jira/browse/SPARK-49294>): Add width attribute for shuffle-write-time checkbox
- [SPARK-49501] (<https://issues.apache.org/jira/browse/SPARK-49501>): Fix double-escaping of table location
- [SPARK-49595] (<https://issues.apache.org/jira/browse/SPARK-49595>): Fix DataFrame.unpivot/melt in Spark Connect Scala Client
- [SPARK-49628] (<https://issues.apache.org/jira/browse/SPARK-49628>): ConstantFolding should copy stateful expression before evaluating
- [SPARK-49695] (<https://issues.apache.org/jira/browse/SPARK-49695>): Postgres fix xor push-down
- [SPARK-49699] (<https://issues.apache.org/jira/browse/SPARK-49699>): Disable PruneFilters for streaming workloads
- [SPARK-49743] (<https://issues.apache.org/jira/browse/SPARK-49743>): OptimizeCsvJsonExpr should not change schema fields when pruning GetArrayStructFields
- [SPARK-49782] (<https://issues.apache.org/jira/browse/SPARK-49782>): ResolveDataFrameDropColumns rule resolves UnresolvedAttribute with child output

- [SPARK-49791] (<https://issues.apache.org/jira/browse/SPARK-49791>): Make DelegatingCatalogExtension more extendable
- [SPARK-49804] (<https://issues.apache.org/jira/browse/SPARK-49804>): Fix to use the exit code of executor container always
- [SPARK-49816] (<https://issues.apache.org/jira/browse/SPARK-49816>): Should only update out-going-ref-count for referenced outer CTE relation
- [SPARK-49819] (<https://issues.apache.org/jira/browse/SPARK-49819>): Disable CollapseProject for correlated subqueries in projection over aggregate correctly
- [SPARK-49829] (<https://issues.apache.org/jira/browse/SPARK-49829>): Fix the bug on the optimization on adding input to state store in stream-stream join
- [SPARK-49836] (<https://issues.apache.org/jira/browse/SPARK-49836>): Fix possibly broken query when window is provided to window/session_window fn
- [SPARK-49843] (<https://issues.apache.org/jira/browse/SPARK-49843>): Fix change comment on char/varchar columns
- [SPARK-49959] (<https://issues.apache.org/jira/browse/SPARK-49959>): Fix ColumnarArray.copy() to read nulls from the correct offset
- [SPARK-49979] (<https://issues.apache.org/jira/browse/SPARK-49979>): Fix AQE hanging issue when collecting twice on a failed plan
- [SPARK-50021] (<https://issues.apache.org/jira/browse/SPARK-50021>): Fix ApplicationPage to hide App UI links when UI is disabled
- [SPARK-50022] (<https://issues.apache.org/jira/browse/SPARK-50022>): Fix MasterPage to hide App UI links when UI is disabled
- [SPARK-50087] (<https://issues.apache.org/jira/browse/SPARK-50087>): Robust handling of boolean expressions in CASE WHEN for MsSqlServer and future connectors
- [SPARK-50176] (<https://issues.apache.org/jira/browse/SPARK-50176>): Disallow reattaching after the session is closed
- [SPARK-50195] (<https://issues.apache.org/jira/browse/SPARK-50195>): Fix StandaloneRestServer to propagate spark.app.name to SparkSubmit properly
- [SPARK-50210] (<https://issues.apache.org/jira/browse/SPARK-50210>): Fix SparkSubmit to show REST API kill response properly
- [SPARK-50235] (<https://issues.apache.org/jira/browse/SPARK-50235>): Clean up ColumnVector resource after processing all rows in ColumnarToRowExec

- [SPARK-50258] (<https://issues.apache.org/jira/browse/SPARK-50258>): Fix output column order changed issue after AQE optimization
- [SPARK-50312] (<https://issues.apache.org/jira/browse/SPARK-50312>): SparkThriftServer createServer parameter passing error when kerberos is true
- [SPARK-50421] (<https://issues.apache.org/jira/browse/SPARK-50421>): Fix executor related memory config incorrect when multiple resource profiles worked
- [SPARK-50433] (<https://issues.apache.org/jira/browse/SPARK-50433>): Fix configuring log4j2 guide docs for Spark on YARN and UT
- [SPARK-50463] (<https://issues.apache.org/jira/browse/SPARK-50463>): Fix ConstantColumnVector with Columnar to Row conversion
- [SPARK-50483] (<https://issues.apache.org/jira/browse/SPARK-50483>): BlockMissingException should be thrown even if ignoreCorruptFiles is enabled
- [SPARK-50492] (<https://issues.apache.org/jira/browse/SPARK-50492>): Fix java.util.NoSuchElementException when event time column is dropped after dropDuplicatesWithinWatermark
- [SPARK-50498] (<https://issues.apache.org/jira/browse/SPARK-50498>): Avoid unnecessary py4j call in listFunctions
- [SPARK-50505] (<https://issues.apache.org/jira/browse/SPARK-50505>): Fix spark.storage.replication.proactive default value documentation
- [SPARK-50510] (<https://issues.apache.org/jira/browse/SPARK-50510>): Fix sporadic ReattachableExecuteSuite failure
- [SPARK-50545] (<https://issues.apache.org/jira/browse/SPARK-50545>): AccessControlException should be thrown even if ignoreCorruptFiles is enabled

Dependency changes

While being a maintenance release we did still upgrade some dependencies in this release they are:

- [SPARK-50150] (<https://issues.apache.org/jira/browse/SPARK-50150>): Upgrade Jetty to 9.4.56.v20240826
- [SPARK-50316] (<https://issues.apache.org/jira/browse/SPARK-50316>): Upgrade ORC to 1.9.5

You can consult JIRA for the detailed changes (<https://s.apache.org/spark-3.5.4>).

We would like to acknowledge all community members for contributing patches to this release.

Spark News Archive (</news/>)

Latest News

Spark 3.5.5 released (</news/spark-3-5-5-released.html>) (Feb 27, 2025)

Spark 3.5.4 released (</news/spark-3-5-4-released.html>) (Dec 20, 2024)

Spark 3.4.4 released (</news/spark-3-4-4-released.html>) (Oct 27, 2024)

Preview release of Spark 4.0 (</news/spark-4.0.0-preview2.html>) (Sep 26, 2024)

Archive (</news/index.html>)



(<https://www.apache.org/events/current-event.html>)

DOWNLOAD SPARK ([/DOWNLOADS.HTML](/downloads.html))

Built-in Libraries:

SQL and DataFrames (</sql/>)

Spark Streaming (</streaming/>)

MLlib (machine learning) (</mllib/>)

GraphX (graph) (</graphx/>)

Third-Party Projects (</third-party-projects.html>)

Apache Spark, Spark, Apache, the Apache feather logo, and the Apache Spark project logo are either registered trademarks or trademarks of The Apache Software Foundation in the United States and other countries. See guidance on use of Apache Spark trademarks (</trademarks.html>). All other marks mentioned may be trademarks or registered trademarks of their respective owners. Copyright © 2018 The Apache Software Foundation, Licensed under the Apache License, Version 2.0 (<https://www.apache.org/licenses/>).